



A Survey of Monte Carlo Tree Search M...

Action-Conditional Video Prediction u...

Reinforcement learning for robots usi...

Deep Learning for Real-Time Atari Gam...

Long Short-Term Memory

Deep visual foresight for Embodiment Learning with Locally Linear La...

MuJoCo: A physics engine for model-based Deep Residual Learning for Image Reco...

Recurrent Environment Simulators

Imagination-Augmented Agents for Deep Reinforcement Learning

Stochastic Actor-Critic for Deep Reinforcement Learning

Between MDPs and Semi-MDPs: A Framework...

The Arcade Learning Environment: An Eval Environment for Deep Reinforcement Learning

Self-supervised control through Reinforcement Learning with Unsupervi...

Massively Parallel Methods for Deep RL

Prioritized Experience Replay

Reinforcement Learning: An Introduction

Integrated Architectures for Learning...

A Distributional Perspective on Reinforcement Learning

Learning to Predict by the Asynchronous Methods for Deep Trust Region Policy Optimization

Dueling Network Architectures for Deep Reinforcement Learning

Deep Reinforcement Learning with Doubly Stochastic Actor-Critic

Playing Atari with Deep Reinforcement Learning

Mastering Atari, Go, Chess and Shogi ...

Rainbow: Combining Improvements in Deep Reinforcement Learning

Value Prediction Network

Distributed Prioritized Experience Replay

Mastering the game of Go with deep reinforcement learning

Learning Continuous Control Policies with Actor-Critic

From Pixels to Torques: Policy Learning with Model-based Reinforcement Learning

PILCO: A Model-Based and Data-Efficient Policy Learning Framework

Value Iteration Networks

Efficient Selectivity and Backup Operation in Value Iteration

Mastering the game of Go without human data

Bandit Based Monte-Carlo Planning

Rectified Linear Units Improve Restricted Boltzmann Machines

Neural Fitted Q Iteration - First Experiments

Monte-Carlo Go Developments